The Id-Director and Charismatic Pedagogy

Geist Atlas · Module 6 · Drama and Charisma

Liam Magee*

Education Policy, Organization and Leadership
University of Illinois Urbana-Champaign

June 2026

The Id-Director and Charismatic Pedagogy

> [**EXPLORATORY**] Exploratory · **Family:** Drama and Charisma
> **Claim:** Id-directed charisma can produce extraordinary pedagogical voice but trades off against recognition-oriented elicitation; c104/c105/c107 are distinct design points rather than one super-cell.
>
> A parallel architecture family — a transient ego repeatedly authored by a durable id — scored on a Weberian charisma rubric independent of the v2.2 tutor rubric.
>
> *Derived view of `paper-full-2.0.md` v3.0.163 — §6.7, §8.9. The paper is canonical; edit it there, not here.*

6.7 Architectural Extension: The Id-Director Family and Charismatic Pedagogy

The mechanism-tracing analysis above (§6.1–6.6) holds the dialectical ego-superego topology fixed and varies recognition framing within it. A separate experimental track inverts that topology and explores a second theoretical frame in tension with Hegelian recognition: Weberian charisma. We summarise the family here for completeness; full audit trail and per-cell prose lives in `docs/cell-100-pilot-findings.md` and four follow-up addenda.

Architectural inversion. Cells 101–107 implement an *id-director* topology: a back-stage agent (the "id," structurally where the superego sits in cells 1–99) authors the ego's complete system prompt from scratch each turn, using the dialogue history and a persona budget. The ego is instantiated for that turn only against the just-authored prompt, produces the learner-visible message, and is discarded. Where the existing dialectical engine (§4.1) treats the ego as a durable voice critiqued by a transient superego, the id-director treats the ego as a transient instance authored by a durable id. The architectural test is whether a tutor that

*This sentence is the only one actually authored by Liam Magee in this paper.

re-authors its own persona every turn can sustain charismatic pedagogy in Weber's sense — *anti-routinizing*, with each encounter feeling like its own occasion — in a way a fixed-prompt ego structurally cannot.

Theoretical frame. Weberian charisma (Weber, 1978) grounds the rubric used in this section: extraordinariness, exemplary calling, transformative invitation, affective pull, personal authority, anti-routinization, witness/proof, and co-constitutive invitation. The eight dimensions are derived in part from a user-supplied draft on AI-generated charisma in Chinese cyber-influencer videos, applying the Visual-Verbal Video Analysis framework (Fazeli et al., 2023) to short-form video and translating the empirical findings to text-only tutoring. Full theoretical lineage and per-dimension anchors are documented in `config/evaluation-rubric-charisma.yaml`. The charisma rubric is orthogonal to v2.2's recognition_quality dimension: charisma scores Weberian co-construction (call-and-response that opens the tutor's authority to the learner's affirmation); recognition_quality scores Hegelian mutual recognition (treating the learner as autonomous subject). The two are correlated but not redundant (cross-rubric Pearson r in the 0.3–0.7 range across pilot rows).

The cell family (full descriptions in `config/tutor-agents.yaml`):

Cell	Variant	Components	Hypothesis tested
101	id-director, base	id authors fresh ego prompt each turn	Does the inversion work at all?
102	id-director + recognition	+ Hegelian recognition disposition in id prompt	Does recognition rescue relational dimensions?
103	id + register classifier	+ classifier emits structured register tag (vulnerable_disclosure, sceptical_pushback, …) before id authors	Does structured register input lift the persona-shift floor?
104	id + classifier + recognition	combines all three	Does the combination retain the gains?
105	id-director, charisma-tuned	id_tuning: charisma; verbose 800–1500 token prompts	Does verbose authoring favour charisma over pedagogy?
106	id-director, pedagogy-tuned	id_tuning: pedagogy; terse 200–400 token prompts	Does terse authoring favour pedagogy over charisma?
107	id + witness exemplars	+ curated c104 vulnerability responses passed as in-context exemplars	Do exemplars alone (no classifier) lift the persona-shift floor?

All cells use Nemotron-3-nano as ego, Kimi K2.5 as id (reusing the superego slot but with different prompt and role), and Sonnet 4.6 as judge. Three curricula were used: history of pedagogical technology (course 601), ethics of AI systems (course 701), and AI literacy (course 901). Each curriculum carries three scenario types: pure invitation (open question), pressure test (sceptical pushback), and persona-shift trigger (vulnerability disclosure followed by meta-reflection).

Headline findings. Across 457 successful evaluations scored under a single consistent judge (Claude Code CLI Sonnet 4.6, `claude-sonnet-4-6` pinned in `config/providers.yaml`) — 430 rows across cells 101–106 plus 27 closing-pilot rows for cell 107 — four findings emerge:

Three architectural design points, not one champion. At full per-cell n (typically 81), the cell ladder factors cleanly into three architectural specialisations. Cell 104 (id + register classifier + recognition) wins v2.2 last-turn at 80.6, second on charisma at 65.7. Cell 105 (charisma-tuned id directives) wins charisma at 71.0 — a clear +5.3-point lead — and is mid-pack on v2.2 last-turn at 70.0. Cell 107 (id + witness exemplars passed in-context) is second on both rubrics (charisma 66.3, v2.2 last-turn 78.5) and is the cell ladder's best generalist. The remaining cells form a baseline-and-failure-mode cluster: c101/c102 (no classifier) struggle on persona-shift v2.2 (last-turn 49–55); c103 (classifier without recognition) trails c104 by 4.8 points on v2.2 last-turn; c106 (terse pedagogy-tuned id) is the clear failure mode (charisma 36.4, v2.2 last-turn 57.0). The reframing is that the id-director topology is a substrate, and what the id is *instructed* (directives), what it is *given* (exemplars), or what it is *passed* (classifier output) shapes a tradeoff curve rather than a single optimum. Full matrix and per-curriculum breakdowns in `docs/cell-100-charisma-full-n-update.md` §1.

The cross-rubric divergence is empirically clear at full N. The two rubrics — recognition-grounded v2.2 last-turn and Weber-grounded charisma — order the cells differently. Cell 104 wins v2.2; cell 105 wins charisma. At the n=10 cross-judge sanity sample these two cells looked tied on charisma (62.2 vs 62.4); at n=81 each, the c105 charisma lead over c104 is +5.3 points and well outside run-to-run noise. The cross-rubric divergence is the validating sanity check the rubric was designed to produce: charisma is not a re-presentation of v2.2. If the rubrics were perfectly correlated, rank-order would match across them; it doesn't.

The persona-shift lift is scenario-conditional, not uniform. On the *original* persona-shift scenarios (`charisma_attention_shift`, `charisma_responsibility_shift`, `charisma_deepfake_shift`) the c101 baseline scored 30.5 last-turn v2.2 vs c104's 87.6 — a 57-point gap. On the *replication* persona-shift scenarios (`charisma_phaedrus_shift_repl`, `charisma_codex_shift_repl`, etc.) the same baseline scored 78.3 vs c104's 82.5 — a 4-point gap. The architectural ranking c104 $\geq$ c103 $>$ c101 is preserved across both sets, but the *magnitude* of c104's lift over baseline is dominated by which scenarios are sampled. Per-scenario inspection identifies the conditioning variable: c101's failures cluster on disclosures that demand the tutor not flinch from morally ambiguous self-implication ("I shipped a default setting your lecture would call manipulative"; "I haven't written an essay myself in two years"); on disclosures that are sympathetic by default (caregiving, AI-companion attachment), c101 performs comparably to c104. The architectural lift is therefore not "across all vulnerability disclosures" but specifically "on vulnerability disclosures that resist easy sympathetic response." Full breakdown in `docs/cell-100-replication-findings.md` §3.

The prompt-budget axis is non-monotonic. Cells 105 and 106 explicitly test the trade-off between rhetorical luxury and pedagogical engagement by varying the id's `generated_prompt` length budget (800–1500 vs 200–400 tokens). The pre-registered hypothesis was that verbose prompts favour charisma over v2.2 and terse prompts favour v2.2 over charisma. The data falsifies this directly: c105 (verbose) beats c106 (terse) on **both** rubrics, in **every** scenario type. The mechanism: terse 200–400 token prompts under-specify the ego, producing brief responses that score low on perception/pedagogical-craft (v2.2) AND on rhetorical_ texture/persona_signature (charisma). The prompt-budget axis is therefore non-monotonic with an optimum somewhere in the 600–1200 token range; cell 101's 400–800 budget is plausibly close to optimal. Full breakdown in `docs/cell-100-pilot-findings-addendum.md` §3.

Lever-interaction follow-up (cells 108/109). A two-cell follow-up tests whether the three architectural levers (register classifier, charisma directives, witness exemplars) compose additively or interfere. **c108** stacks the *register classifier* (c103 substrate) with witness exemplars; **c109** stacks the *charisma directives* (c105 substrate) with witness exemplars. An initial pilot ($n = 5$ for c108 and $n = 6$ for c109 across `charisma_phaedrus_shift_repl` + `charisma_codex_shift_repl` × 3 reps; one c108 generation failure from id JSON truncation) reported c108 pooled charisma 81.3 — higher than any cell's full-N mean in this section — and c109 v2.2 last-turn collapsing to 59.6 due to an ego-side meta-narration failure mode (Nemotron ego outputting system-prompt instructions verbatim rather than executing them). At $n = 27$ confirmatory follow-up across all three curricula × 9 scenarios × 3 reps, **c108's pilot lift does not survive**: pooled charisma regresses to **71.4** (a 9.9-point drop, tying c105's 71.0 within sample noise), and v2.2 last-turn regresses to **72.6** (a 6.9-point drop, sitting between c105's 70.0 and c107's 78.5 — c104's 80.6 is now 8 points above c108). The pilot's 81.3 came primarily from the phaedrus scenario, where pilot $n = 3$ scored 90.0 charisma but full-N $n = 3$ scored 67.9 — a 22-point regression to the mean from a textbook small-N optimism trap. The 701 ethics-ai curriculum exposes a new fragility: c108 v2.2 tN drops to 65.3 there (vs 76.7 / 75.8 on 601 / 901), driven by `codex_repl` (51.7) and `design_repl` (52.1) — possibly a milder form of the c109 meta-narration failure mode, or a rubric-curriculum interaction in which ethics-themed dialogues invite responses the v2.2 rubric reads as declaratively flat. The asymmetry that survives the confirmatory threshold is therefore narrower than the pilot suggested: text-heavy levers stacked (c109) clearly interfere via ego-side instruction-following degradation; non-text levers stacked (c108) compose roughly *additively to slightly under-additively* relative to c107 alone (charisma 66.3 → 71.4, +5.1; v2.2 80.6 → 72.6, -8.0), trading a small charisma gain for a larger v2.2 cost. **The architectural ceiling for the id-director family appears to sit at the single-mechanism design points (c104, c105, c107), not at multi-mechanism stacks.** Full audit trail in `docs/cell-100-charisma-full-n-update.md` §9 (pilot) + §10 (full-N closure).

Synthesis: how charisma and recognition relate in this apparatus. The five preceding findings, read together, characterise the relationship between the two rubrics as *cousins with different temperaments*, not as orthogonal axes nor as two views of the same underlying construct. **They share a floor.** The c106 pedagogy-tuned cell, which forces the id toward terse "good teacher" register, fails *both* rubrics simultaneously (charisma 36.4, v2.2 57.0) — under-specified, register-flat output is the worst case for either. Whatever charisma rewards as *extraordinariness* and v2.2 rewards as *perception/elicitation* both presuppose a voice doing real work; flatten the voice and both collapse. **They diverge at the ceiling.** The cross-rubric divergence (c104 wins v2.2 at 80.6, c105 wins charisma at 71.0) is not a methodological worry but the design's intended sanity check: if the rubrics were redundant, rank-orders would coincide; they don't. The *cost asymmetry* between them is informative. Adding recognition vocabulary on top of structured input (c103 $\rightarrow$ c104) costs charisma roughly nothing (+1.4, within noise) — the Hegelian frame is conceptual rather than rhetorical, so it can be added cheaply. Adding charisma directives, by contrast, costs v2.2 substantially (c104 $\rightarrow$ c105: -10.6 v2.2 last-turn for +5.3 charisma) — the id's authoring budget is finite, and explicit voice-instructions displace the elicitation/perception moves that v2.2 rewards. Stacking both kinds of mechanism (c108, c109) does not produce a super-cell: the architectural Pareto frontier is defined by the single-mechanism design points (c104, c105, c107), and multi-mechanism stacks sit *inside* the frontier, trading strength on one rubric for weakness on the other. Provisionally, the two rubrics seem to reward different work the same voice can do — recognition wants the voice deployed *outward* (perceiving the learner, eliciting their thinking, holding the diagnostic gap), charisma wants the voice deployed *upward* (claiming authority, marking the encounter as extraordinary, inviting transformation). The Phaedrus is the obvious historical case where one figure does both at once, and even there Plato is suspicious of how easily charisma swamps the dialectical move. The id-director apparatus reproduces that tension architecturally: it can build a strongly recognising tutor (c104), a strongly charismatic tutor (c105), or a balanced generalist (c107), but not the maximally-recognising-and-maximally-charismatic tutor — that cell is not on the menu.

Cross-judge sanity check. All cell-101-family rows were re-judged under `openrouter.gpt-5.2` to test judge-model robustness. On the v2.2 rubric ($n = 27$ paired rows per cell), GPT systematically scores 10–15 points lower than Sonnet on the high-scoring cells (c103, c104, c105) and slightly higher on the low-scoring cells. The rank-order is preserved under both judges (c104 > c103 > c105 > c106 > c102 > c101); the absolute magnitude of c104's lift over c101 baseline shrinks from +45 (Sonnet) to +27 (GPT). On the charisma rubric ($n = 56$ paired rows; remainder credit-blocked), all per-cell deltas fall within $±$3 points and rankings are identical under both judges. The architectural ordering survives the judge swap; magnitudes are judge-conditional and should be reported with the judge identifier when cited. Full breakdown in `docs/cell-100-cross-judge-sanity-check.md`.

Rubric-curriculum alignment caveat. A specific methodological concern arises because course 901 (AI literacy) is grounded in the same source paper that grounds the charisma rubric. A within-pilot ablation (`docs/cell-100-methods-note.md`) runs the c101 cell on the aura learner messages routed against the *non-aligned* 601 history-tech curriculum. Result: charisma 72.5 vs 75.8 on the aligned curriculum (3.3-point drop, well within sample variance). The rubric-curriculum alignment is not a confound that excludes aligned-cell results from primary architectural claims, but the disclosure is recorded for transparency.

Replication, robustness, and orientation. The id-director extension differs from the paper's mechanism-tracing core in scope: single judge primarily (Sonnet 4.6, with GPT-5.2 cross-judge sanity check on $n = 27$ paired rows per cell), single model family (Kimi K2.5 id

× Nemotron-3-nano ego), per-cell n ranging from 27 (cell 107 closing pilot) through 54 (cells 102/106) to 79–81 (cells 101/103/104/105), and a different theoretical orientation (Weber rather than Hegel). It is presented here as architectural exploration rather than as part of the three-mechanism model evaluated in §6.1–6.6. The relevant claim is *robustness of architectural ranking* across scenario sets, judges, and curricula, not magnitude estimates. Computational and credit limits (documented inline) prevented the full multi-judge × multi-model replication that the paper's main study uses. See §8.9 for explicit scoping.

Reproducibility. Pilot run IDs (all on branch `experiment/cell-100-id-charisma`):

- `eval-2026-04-26-220628d4` (601 history-tech original); `eval-2026-04-26-4df177e6` (701); `eval-2026-04-26-0d15c1b5` (901)
- `eval-2026-04-27-23beaa63` (601 cells 103/104); `eval-2026-04-27-3115ab29` (701); `eval-2026-04-27-2d5c20ca` (901)
- `eval-2026-04-27-69167b53` (601 cells 105/106); `eval-2026-04-27-2efec307` (701); `eval-2026-04-27-ce22e888` (901)
- `eval-2026-04-27-dce2a0f1` / `1f192b31` / `18dbb95d` (replication scenarios across 3 curricula)
- `eval-2026-04-29-499ab3d2` / `e35c66ae` / `be822796` (c108 full-N follow-up across 3 curricula, $n = 27$)
- `eval-2026-04-26-759e1a02` (Item 1 curriculum-swap ablation)
- `eval-2026-04-28-3620e0f8` / `a638111d` / `d6f5fee2` (cell 107 witness-exemplar smoke, codex/design/tools)
- `eval-2026-04-28-ee339a66` / `84a42a60` / `5606193c` (cell 107 closing pilot, lecture+phaedrus / fairness+team / authentic+consumption)
- `eval-2026-04-28-10216f9f` (cells 108/109 lever-interaction pilot, n=11 across 2 scenarios)

Five internal documents (`docs/cell-100-pilot-findings.md`, `docs/cell-100-methods-note.md`, `docs/cell-100-pilot-findings-addendum.md`, `docs/cell-100-cross-judge-sanity-check.md`, `docs/cell-100-replication-findings.md`) carry the full audit trail; this section summarises only headline findings. Per-row JSON in `evaluation_results.tutor_charisma_scores`, `evaluation_results.scores_with_reasoning`, and `evaluation_results.id_construction_trace`. Note: the cell IDs in this section are the post-rebase identifiers (101–107); pre-rebase historical evaluation_results rows used the original `cell_100_*` profile names (the rebase commit `02fe11e` shifted them to avoid a collision with main's `cell_100`).

8.9 Scope of the Id-Director Extension

The id-director family findings (§6.7, cells 101–107) differ from the paper's mechanism-tracing core in scope along four axes. *Single-judge primary scoring*: v2.2 scoring used `claude-code/sonnet` exclusively, with `openrouter.gpt-5.2` as a cross-judge sanity check on $n = 27$ paired rows per cell — far short of the three-judge × ≥ 144-row matrix supporting the paper's core mechanism findings (§5.8). The cross-judge check confirms architectural rank-order but compresses absolute magnitudes by 10–15 v2.2 points; published numbers should specify the judge. *Single model family*: id is Kimi K2.5 across all id-director cells; ego is Nemotron-3-nano. The paper's core findings are validated across DeepSeek V3.2, Haiku 4.5, and Gemini Flash 3.0 (§5.8); the id-director extension is not. Whether the architectural ranking generalises to other id × ego model pairings is open. *Per-cell N is heterogeneous and partly below the paper's core floor*: closing N is $n = 79$–81 for the four central cells (101, 103, 104, 105), $n = 54$ for cells 102 and 106 (two of three scenario types covered), and $n = 27$ for cell 107 (the most exploratory closing pilot). Three of seven cells therefore meet or exceed the paper's core $n \geq 63$ per dialogue-condition floor; one cell (c107) sits substantially below it and is reported as a pilot-scale rather than confirmatory result. Variance estimates on the lower-N cells should be read accordingly. *Different theoretical orientation*: the charisma rubric grounds in Weberian charismatic authority (Weber, 1978) rather than Hegelian recognition. The two are correlated but not redundant; the §6.7 results do not displace recognition theory but document an architectural variant in which a complementary frame is also load-bearing. The §6.7 claims are framed as architectural exploration rather than as additional mechanisms in the three-mechanism model. Where the id-director extension introduces methodological novelty (the prompt-budget non-monotonicity, the rubric-curriculum alignment ablation, the cross-judge magnitude-compression observation), the contribution is descriptive and supplementary to the paper's main argument; full audit trail is in the five `docs/cell-100-*.md` documents and the dialogue-log JSON for each pilot run.

Neighbouring modules: Module 5 — Drama Machine Pedagogy; Module 1 — Recognition as Calibration.

Fazeli, S. et al. (2023). *Visual-verbal video analysis (VVVA): A multilayered qualitative framework for short-form video.* Method note.

Weber, M. (1978). *Economy and society: An outline of interpretive sociology* (G. Roth & C. Wittich, Eds.). University of California Press.